

An Introduction to Data-Oriented Programming

Yaser Zhian

Dead Mage Studio

Target Audience

Programmers who face:

- Hard problems
- Performance/resource constraints
- Many solutions

Agenda

- Define a simple sample problem
- Discuss solutions
- Go over couple of solutions (hopefully today)
- Problem areas and reasoning about behavior
- Another (better) solution
- Theoretical groundwork for “data-oriented” programming

Problem Statement

- Simulate the behavior of, and display many ants crawling on a predefined map, from “city” to “city”.
 - Various behaviors
 - Each ant is a pixel
 - This is a benchmark; we’re going to evaluate multiple solutions
 - Not particularly focused on performance
- What more do you need to know?

Let's Discuss

- What is our map?
- What's an ant? What can it do?
 - “Ant” stories?!
- How do we render?
- How to organize the solution?
 - Make it work,
 - Make it right,
 - Make it fast.
- What else?

Coding

Part 1

Where Is My Time?!

- How long should it take?
 - Let's count...
 - How many operations do we have? (ALU/FPU*, branches, memory, etc.)
 - How long do each of these take?
- Cost of abstractions
- ALU/SSE operations
 - Throughput vs. Latency
- Branch (mis)predictions
- Memory

Memory

- What do we mean by “speed”?!
- What does the spec say?
- How can we measure?
- Let’s measure...

Coding

Part 2

RAM isn't RAM (or, why is RAM so slow?!)

- The “stored program” machine
- RAM is S...L...O...W...
- RAM isn't random-access
- Other problems
 - Multiple cores
 - NUMA
 - Mispredicted branches in a slow-RAM world
- What to do?

Cache

- General operation (levels, algorithms)
- Cache line
- Associativity
- Multi-CPU/Multi-Core
- “Why” does a cache work?!!!

Costs (on a Haswell)

- (Mostly throughput)
- ALU:
 - Mostly in 0.5-2 cycle range
 - 64-bit (signed) division is ~20-80 cy
- Branches: mostly in 1-2 cy range
- FPU/SSE/AVX/etc.:
 - Mostly in 0.5-2 cy range
 - (Various incarnations of) division and sqrt are in 7-28 cy range
 - Transcendentals (sin, cos, tan, etc.) are in ~50-150 cy range
- Cost of hitting main memory: ~100-200 cycles

Optimization

- If “ p ” is the portion of time that we spend in processing, and we improve that by a factor of “ s ”, then the “speed up” we gain is:
 - $speedup = \frac{1}{(1-p) + \frac{p}{s}}$
 - This is called “Amdahl’s Law”
- One result of the above law is this:
 - It’s always better to optimize the larger (largest) parts!
 - It’s all about memory access
- So let’s write some code...

Coding

Part 3

Data-Oriented Design Principles

Or “Gospel according to Mike (Acton)”

Step 0:

“The purpose of all programs, and all parts of those programs, is to transform data from one form to another.”

Step 1:

“If you don’t understand the data you don’t understand the problem.”

“Conversely, understand the problem by understanding the data.”

Step 2:

“Different problems require different solutions.”

Step 3:

“If you have different data, you have a different problem.”

Step 4:

“If you don’t understand the cost of solving the problem, you don’t understand the problem.”

Step 5:

“If you don’t understand the hardware,
you can’t reason about the cost of solving
the problem.”

(also, understand your code; i.e. language, compiler, OS, etc.)

Step 6:

“Everything a data problem. Including usability, maintenance, debug-ability, etc. Everything.”

Step 7:

“Solving problems that you probably don’t have, creates problems that you definitely do.”

Step 8:

“Latency and throughput are only the same in sequential systems.”

Step 9:

“Software does not run in a magic fairy aether powered by the fevered dreams of CS PhDs.”

Step 10:

“Reason must prevail.”

- This page is intentionally left blank!

Three Big Lies (#1)

Software is a platform

Three Big Lies (#2)

Code should be designed
around a model of the
world

Three Big Lies (#3)

Code is more important
than data

More Lies

- If you put in enough abstractions, the solution becomes self-evident.
- You can abstract the real world away.
- Ignorance is bliss.
- Most misused quote:
 - “Premature optimization is the root of all evil.”

Some Final Thoughts

- Cult of Ignorance
 - Reusability
 - Patterns!
 - Algorithms vs. patterns
 - Genericity
 - Performance
- Encapsulation for maintainability vs. data hiding
- Abstraction vs. transformation
- Ignoring the real world doesn't make it disappear!
- Top-down vs. bottom-up vs. side-to-side

Beware of delusion!

- Everything can be used to justify anything.
- There is no “One True Way”.
- Look at the data, and write the simplest solution that satisfies all constraints.
 - Input/output data
 - Meta data: how many, how often, when
 - Constraints: time, ability, cognitive load, hardware

Thank You!

- References

- CppCon 2014: Mike Acton, “Data-Oriented Design and C++”
<https://www.youtube.com/watch?v=rX0ItVEVjHc>
- Richard Fabian, “Data-Oriented Design”
<http://www.dataorienteddesign.com/dodbook/>
- Agner Fog, Software Optimization Manuals (Microarch and timings)
<https://www.agner.org/optimize/>

- Any questions?